

**SFWRENG 4F03**  
**Parallel Computing**  
**Winter 2020**

## **02 Parallel Computing**

Dr Asghar Bokhari

Faculty of Engineering, McMaster University

January 6, 2020



# How to handle a large complex task?

- Two ways:
  - ▶ Sequential: One person works on different parts of the task one after the other
  - ▶ Parallel: A number of persons work on different parts (in parallel) that are combined to complete the task.

Examples:

- ▶ Move all chairs of this room to another room
- ▶ Construction - different tasks by specialised workers
- ▶ Making cars - different sub-assemblies by different companies

# How to handle a large complex task?

- Advantage of second approach:
  - ▶ Same job in less time
  - ▶ Much bigger jobs can be handled in reasonable time


Steps involved:

- ▶ Decomposition of task into smaller independent parts
- ▶ Employing many workers each working on the smaller part
- ▶ Coordination

# Jig saw puzzle example <sup>1</sup>



## Serial Computing



Suppose you want to do a jigsaw puzzle that has, say, a thousand pieces.

We can imagine that it'll take you a certain amount of time. Let's say that you can put the puzzle together in an hour.



Supercomputing in Plain English: Overview  
Wednesday August 29 2007

---

<sup>1</sup>Taken from OU Supercomputing Center for Education and Research,  
University of Oklahoma

# Jig saw puzzle example



## Shared Memory Parallelism



If Horst sits across the table from you, then he can work on his half of the puzzle and you can work on yours. Once in a while, you'll both reach into the pile of pieces at the same time (you'll contend for the same resource), which will cause a little bit of slowdown. And from time to time you'll have to work together (communicate) at the interface between his half and yours. The speedup will be nearly 2-to-1: y'all might take 35 minutes instead of 30.




Supercomputing in Plain English: Overview  
Wednesday August 29 2007

# Jig saw puzzle example



## The More the Merrier?




Now let's put Bruce and Dee on the other two sides of the table. Each of you can work on a part of the puzzle, but there'll be a lot more contention for the shared resource (the pile of puzzle pieces) and a lot more communication at the interfaces. So y'all will get noticeably less than a 4-to-1 speedup, but you'll still have an improvement, maybe something like 3-to-1: the four of you can get it done in 20 minutes instead of an hour.



Supercomputing in Plain English: Overview  
Wednesday August 29 2007

# Jig saw puzzle example



## Diminishing Returns

If we now put Rebecca and Jen and Alisa and Darlene on the corners of the table, there's going to be a whole lot of contention for the shared resource, and a lot of communication at the many interfaces. So the speedup y'all get will be much less than we'd like; you'll be lucky to get 5-to-1.

So we can see that adding more and more workers onto a shared resource is eventually going to have a diminishing return.



Supercomputing in Plain English: Overview  
Wednesday August 29 2007

# Jig saw puzzle example



## Distributed Parallelism



Now let's try something a little different. Let's set up two tables, and let's put you at one of them and Horst at the other. Let's put half of the puzzle pieces on your table and the other half of the pieces on Horst's. Now y'all can work completely independently, without any contention for a shared resource. **BUT**, the cost of communicating is **MUCH** higher (you have to scootch your tables together), and you need the ability to split up (*decompose*) the puzzle pieces reasonably evenly, which may be tricky to do for some puzzles.



Supercomputing in Plain English: Overview  
Wednesday August 29 2007



# Jig saw puzzle example



## More Distributed Processors




It's a lot easier to add more processors in distributed parallelism. But, you always have to be aware of the need to decompose the problem and to communicate between the processors. Also, as you add more processors, it may be harder to load balance the amount of work that each processor gets.



Supercomputing in Plain English: Overview  
Wednesday August 29 2007

# Jig saw puzzle example




**Load balancing** means giving everyone roughly the same amount of work to do.

For example, if the jigsaw puzzle is half grass and half sky, then you can do the grass and Julie can do the sky, and then y'all only have to communicate at the horizon – and the amount of work that each of you does on your own is roughly equal. So you'll get pretty good speedup.



Supercomputing in Plain English: Overview  
Wednesday August 29 2007

# Jig saw puzzle example



## Load Balancing

Load balancing can be easy, if the problem splits up into chunks of roughly equal size, with one chunk per processor. Or load balancing can be very hard.



Supercomputing in Plain English: Overview  
Wednesday August 29 2007

# Parallel Programming

- Traditionally computers carry out tasks by executing one instruction at a time sequentially
- We can solve bigger problems in less time by using more than one computers in parallel.
- But
  - ▶ Processing power of CPUs have increased tremendously recently
  - ▶ Some of these CPUs can execute one instruction in 9 to 12 billionth of a second
  - ▶ Is that not fast enough?
- Several problems require much more speed
- Many problems interesting to scinetists and engineers will take months, even years to run on a PC but may take a few hours on a supercomputer.

# Parallel Programming

- There are problems that may not fit on a PC because they need more RAM or disk space. Some examples of these problems:
  - ▶ Simulation
  - ▶ Climate modeling
  - ▶ Fluid Turbulence
  - ▶ Pollution Dispersion
  - ▶ Human Genome
  - ▶ Ocean Circulation
  - ▶ Galaxy information
  - ▶ Data Mining
  - ▶ Visualization - turning large amounts of data into pictures that scientists can understand.

# Parallel Programming

- But, I am not into all these things!!
  - ▶ Learning how to run your code faster, empowers you to handle more complex problems in less time
  - ▶ Historically, what happens in supercomputing today, will be in your desktop 10 to 15 years from now
  - ▶ Learning it today puts you ahead of the curve.
  - ▶ Applications such as www, web services, on line banking and on line retailers, oil exploration, data mining, advanced graphics and virtual reality, computer aided diagnostics in medicine already use clusters to handle complexity and traffic.
  - ▶ parallelism is not restricted to supercomputers alone but it is getting popular in workstations and personal computers also. Future programs will be required to exploit multiple processors inside each computers in addition to those available over the network

# Alternatives

- Traditional Supercomputers
  - ▶ Very fast single computers with large memory bandwidth
  - ▶ Some successful applications
  - ▶ Use sequential programming that is most commonly understood by programmers
  - ▶ Supported by a lot of experience in compilers and other tools
  - ▶ Interconnects not required
  - ▶ very expensive
  - ▶ generates a lot of heat requiring special cooling
  - ▶ performance of a single CPU reaching its limit

# Alternatives

- **Use more than one computers working in parallel**
  - ▶ Take advantage of the advances made in processor, networks and memory fields
  - ▶ Use off-the-shelf components
  - ▶ Possibilities to quickly integrate new elements into a system for scalability or for taking advantage of improvements in a particular technology
  - ▶ Usually much cheaper compared to single processor supercomputers
  - ▶ Sequential programs won't run
  - ▶ Not many programmers know parallel programming techniques
  - ▶ Mature tools and compilers not available



# Alternatives

- A number of different arrangements have been tried
- Location of processors
  - ▶ Many processors (CPUs) located inside a single enclosure
  - ▶ A number of computers interconnected together to act as a supercomputing platform - multicomputer systems, clustering
- Interconnection
  - ▶ Different types of direct connections
  - ▶ Connection via Ethernet
- Access to memory
  - ▶ Shared memory
  - ▶ Distributed memory
  - ▶ Distributed shared memory
  - ▶ Use of cache memory

# Alternatives

- Software considerations
  - ▶ How to decompose a program into separate independent parts that can run in parallel on a single or multiple processes?
  - ▶ Parallel algorithms
  - ▶ Operating systems to handle parallel programs

# Alternatives

- Different parallel programming models
  - ▶ Shared memory: tasks share a common address space to read from and write to, asynchronously.
    - ▶ Access to shared memory controlled by locks, semaphores and other similar mechanisms
    - ▶ Program development may be simplified
    - ▶ Compilers translate user program variables to actual global memory addresses
    - ▶ Management and understanding of data locality difficult
  - ▶ Message passing
  - ▶ Instruction level parallelism
  - ▶ Multithreading
  - ▶ Data parallel systems
  - ▶ Hybrid systems

# Road Map

- Review of single processor architecture
- CISC, RISC
- Memory and cache organization
- Details of different hardware arrangements/classifications for parallelism
- Software considerations
- Message Passing Computing
- Message Passing Interface (MPI)
- Parallel Virtual Machine (PVM)
- Pthreads
- OpenMP
- Strategies and algorithms
- Applications

# Title

- 
- 
- Template!