

SFWRENG 4F03
Parallel Computing
Winter 2020

06 Parallel Software

Dr Asghar Bokhari

Faculty of Engineering, McMaster University

February 6, 2020



Types of Parallelism

- Parallelism in Hardware

- ▶ Uniprocessor
 - ▶ Pipelining
 - ▶ Multiple issue processor - Superscaler
- ▶ Multiprocessors
 - ▶ SIMD - Vector Processors, GPUs
 - ▶ Shared memory multiprocessors
 - ▶ Distributed memory multiprocessors
 - ▶ Computer clusters

- Parallelism in Software

- ▶ Very little commodity software making use of parallel hardware
- ▶ Software engineers must learn how to write applications that exploit the available hardware.

Terminology

- Threads
 - ▶ Shared memory systems start a process and fork several threads.
 - ▶ The problem is divided among a number of threads.
- Processes
 - ▶ The problem is divided among processes in distributed memory systems
- When the discussion applies equally well to both types of systems process/thread may be used

SPMD

- We start discussing software for MIMD systems
- Instead of creating different programs to be run on each processor, it is possible to write a single program
- SPMD (single program multiple data)
- A single executable behaves as if it were multiple different programs.

Example:

```
if( I am thread/process 0)
    do task1
else
    do task 2
```

- Task parallelism

SPMD

- SPMD programs can implement data parallelism also

Example:

```
if( I am thread/process 0)
    Operate on first half of the array
else
    Operate on the second half
```

Parallel Program Design

- How to design parallel applications?
- Have a serial program and want to parallelize it
- No silver bullet
- A problem may have several parallel solutions
- Different strategies and tools
- Goal - to be able choose among the many available methods to design efficient parallel applications

- Four stage methodology - PCAM (Ian Foster)¹

- 1: Partition

- ▶ Divide the problem into smaller tasks
 - ▶ Fine grained decomposition (sand is easier to pour than bricks)
 - ▶ Focus on the opportunities for parallelism ignoring the number of processors on target computer system
 - ▶ First decompose data associated with problem to approximately equal sized chunks
 - ▶ Associate each operation with the data on which it operates

¹Designing and Building Parallel Programs, by Ian Foster

<http://www-unix.mcs.anl.gov/~itf/dbpp/>

- Four stage methodology - PCAM (Ian Foster)²


- 1: Partition Continued

- ▶ This results in a set of tasks with some data and a set of operations on it
 - ▶ Representing each box as a task (ref next slide)-
Fine-grained decomposition: large number of small tasks
 - ▶ Reoresenting a row as a task - Coarse-grained
decomposition: small number of large tasks


²Designing and Building Parallel Programs, by Ian Foster

<http://www-unix.mcs.anl.gov/~itf/dbpp/>

Partitioning



(a) Square region for each process



(b) Row region for each process

2: Communication

- ▶ An operation may require data from other tasks
- ▶ Data must be transferred between tasks so that computation may proceed
- ▶ Need to have communication strategies - point-to-point or broadcast

3: Aggregation (or Agglomeration)

- ▶ Re-evaluate the task and communication structures defined in above two stages
- ▶ Combine the tasks into larger tasks to improve performance or to reduce costs


4: Mapping

- ▶ Assign tasks to the available processors with a goal to maximize process utilization and minimize communication costs

Task Dependency Graph

- Task: programmer-defined unit of computation
- A node (circle or ellipse) represents a task in the graph
- Edge represents control dependence
- Start Node: node with no incoming edge
- Finish Node: node with no outgoing edge


Dependency Graph



Degrees of Concurrency


- **Maximum degree of concurrency:** the maximum number of tasks that can be executed in parallel at any stage of execution
- **Average degree of concurrency:** the average number of tasks that can be executed in parallel
- The average degree of concurrency is a more useful measure
- **Critical path:** the longest directed path between any pair of start and finish nodes
- **Critical path length:** sum of the weights of the nodes on a critical path
- Average degree of concurrency = $\frac{\text{total amount of work}}{\text{critical path length}}$
- Weight: units of time required to complete a task
- Speedup = sequential execution time / parallel execution time

Dependency Graph




- maximum degree of concurrency is 6
- critical path length is 5
- total amount of work is 14 (assuming each task takes one unit of time)
- average degree of concurrency is $14/5=2.8$
- max achievable speedup = $14/5 = 2.8$

Assignment to 2 processes




- First number is the process number
- Number of tasks per process = $14/2 = 7$ (use ceiling if fraction)
- critical path length is 5
- speedup = $14/7 = 2$
- average degree of concurrency is $14/5 = 2.8$

Assignment to 3 processes



- First number is the process number
- critical path length is 5
- speedup = $14/5 = 2.8$

Assignment to 4 processes



- First number is the process number
- critical path length is 5
- speedup = $14/5 = 2.8$ What is the speedup for 8 processors??

Example



Title

-
-
- Template!